

# Beyond Prompts: How to Build Real AI Products

---



# Agenda



## GenAI Fundamentals

Market landscape and valuations  
LLMs, tokens, pricing  
Training stack & capabilities



## Products & Ops

Agents, tools, orchestration  
Low-code builders, MCP  
Toolkits & operating practices



## Inside AIPM

Context engineering, RAG  
Prompt strategies, evaluation  
Fine-tuning & decision flows



## Also in this session

Case study: Granola  
Key insights • Glossary • Q&A

Example: Notebook LM



Section I

# GenAI Fundamentals

---

What's driving adoption, how LLMs work, training stack, and core capabilities



# GenAI Market: Scale & Momentum

Valuations and growth since 2022 (illustrative figures from session)

OpenAI

**\$500B**

Valuation



Anthropic

**\$183B**

Valuation



xAI

**\$200B**

Valuation



Microsoft

**\$1.8T** → **\$3.8T**

Valuation growth



Google

**2.5x**

Growth



Meta

**\$300B** → **\$1.8T**

Valuation growth



Startups:

Perplexity \$15B

Glean \$7B

11 Labs \$6B

Nvidia ~10x growth

Why now: proven daily utility vs. prior hype cycles (e.g., blockchain/NFT).




# Next-Token Prediction

- Given prior tokens (context), the model predicts the most likely next token, repeatedly forming text.
- Tokens are numeric units of text; roughly 2 words  $\approx$  3 tokens.
- Outputs are probabilistic (sampling from a distribution), so the same prompt can yield different answers.

GPT-4 vocab  $\sim$ 100k tokens

Pricing  $\approx$  \$1.25 / 1M tokens





# Tokens & Pricing Essentials

Key metrics and practical cost-control tips

Unit

**Tokens (input + output)**

Rule of thumb

**~2 words ≈ 3 tokens**

Vocabulary size

**GPT-4 ~100k tokens**



Typical pricing (model-dependent)

**≈ \$1.25 per 1M tokens**

Billable = input + output

Track usage by endpoint



**Practical tips**

Optimize prompts (be concise; specify format)

Use RAG to shrink context to only relevant chunks


Cache & reuse results; avoid repeated calls

Measure token usage per feature to guide optimization.



# Training Stack: Pre-training → Fine-tuning → RLHF

Three stages to build helpful, safe, and capable LLMs



Transformers (2017 “Attention Is All You Need”) enable parallel attention over tokens.



GPUs provide the parallel compute required for training and inference at scale.



# Compute & Output Nature


GPUs enable parallelism • LLM outputs are probabilistic



## Compute: GPUs & Parallelism

- GPUs run massive parallel math, accelerating attention operations in transformers.
- Enables feasible training/inference at scale with large token contexts.
- Throughput and latency depend on model size, batching, and hardware.

Parallel lanes (illustrative)



## Outputs: Probabilistic Sampling

- LLMs return a distribution over the next token; responses may vary for the same prompt.
- Control behavior using temperature, top-p, and system prompts.
- For reproducibility, fix seeds where supported and constrain output formats.

Next-token probabilities (illustrative)



Temperature

Top-p

System prompt



# What LLMs Can Do

Three core capabilities that power intelligent experiences



## Understand

- Interpret intent and context
- Classify, extract entities, detect sentiment
- Ground responses with provided context



## Transform


- Summarize, rewrite, translate
- Answer questions over documents (RAG)
- Generate or refactor code with constraints



## Generate

- Create net-new text, outlines, and drafts
- Produce images or audio via connected tools
- Synthesize plans, ideas, and code snippets

Shift: from CRUD systems to intelligent workflows with a reasoning “brain”. Combine capabilities for end-to-end experiences.



Section II

# Inside AIPM

From context engineering and RAG to prompts, fine-tuning, and decision frameworks



# Case Study: Granola (AI Meeting Assistant)

User value • Business value • Design value • Key features



## User Value

- Saves time with automatic notes and summaries
- Improves accuracy vs. manual note-taking
- Lets participants focus on the discussion



## Business Value

- Large productivity market (meetings at scale)
- Strategic fit with Zoom/Google Meet ecosystems
- Recurring value drives retention and upsell



## Design Value

- Minimal friction: auto-join, seamless transcript
- Clear summaries with action items
- Privacy controls and easy sharing




## Key Features (Engineering)

- Audio capture via meeting APIs or bots
- Speech-to-Text (Whisper/Google)
- LLM system prompts for summaries + action items
- Growth: bot visibility and post-meeting email recaps

Start from user and business value, then design for minimal friction; implement the simplest reliable pipeline first.

# Granola Architecture — Audio → STT → LLM Summary

Simple, reliable pipeline powering meeting notes and growth



## Growth Tactics

Bot presence in meetings → viral loop

Post-meeting email recaps

Shareable links & collaboration

Lightweight onboarding & privacy controls

Note: Obtain consent, provide opt-outs, and secure storage to build trust.



Context Engineering


# Principles

- Respect token limits — include only the most relevant context.
- Decompose tasks — structure instructions, steps, and constraints clearly.
- Prefer retrieval (RAG) over dumping large documents into the prompt.


Token budget awareness

Relevance > Volume

## Context Window (tokens) — Curate the prompt



## Retrieval-first flow




Retrieve only what's relevant; cite sources/IDs for traceability.



# RAG Pipeline — End-to-End (7 Steps)

Chunk → Embed → Store → Query → Retrieve → Prompt → Generate



Chunking: semantic, ~200–500 tokens

Vector DB: Pinecone, Supabase

Include source IDs for traceability

# Prompt Engineering — Core Strategies

Five patterns to increase control, consistency, and quality



## Role assignment

Set persona, seniority, and domain expertise.

Act as a senior product manager



## Multistep reasoning

Ask to think step-by-step or plan before answering.

Think step-by-step; list assumptions first



## Constraints

Define boundaries: cite sources, avoid speculation, and stick to provided context.

Do not assume facts not in the context



## Output format

Specify JSON, tables, or bullet structure explicitly.

Return JSON: {"summary": "...", "actions": []}



## Few-shot prompting

Provide 1–3 examples to guide style and format.

Example → Response pairs

# Prompt Engineering — Examples & Best Practices

Code-style prompts with operational tips and recommended tools



## Example — Structured Summary

```
Act as a senior PM. Summarize the meeting in 5 bullets.  
Return JSON with 3 action_items [{ "title", "owner", "due_date" }].  
// Constraints: use only transcript; no assumptions
```



## Best Practices

- Be explicit: role, output format, constraints.
- Induce reasoning: ask to plan or think step-by-step.
- Use few-shot examples to shape style and structure.
- Version prompts and evaluate changes before rollout.

[Prompt directory](#)[Versioning](#)[Evaluation](#)

## Recommended Tools

### Prompt Wildcard

Reusable prompt variables and templates

### Anthropic Prompt Creator

Iterate prompts with built-in guidance



## Example — Guardrails

```
Use only provided context. Cite source_ids for facts.  
If missing info, reply: "Insufficient context".
```

Note: Track prompt changes alongside model versions and context configs.



# Fine-tuning — Full vs PEFT (LoRA/Adapters)

Adapt base LLMs for domain and style with clear cost/latency tradeoffs



## Full Fine-tuning



- Highest task performance ceiling
- Deep domain adaptation and control
- Best for complex, specialized skills



- Expensive compute and longer training
- More data required; careful safety tuning
- Higher latency/serving cost for large models

All params updated



Example: BloombergGPT (finance)



✓ Use when: deep domain accuracy is critical



## PEFT — LoRA/Adapters



- Much cheaper and faster to train
- Lower latency and easier deployment
- Swap adapters per domain/persona

Small subset updated



- May not match full fine-tune peak performance
- Limited for drastic behavior shifts
- Still requires quality labeled data



Example: Bank AI (tone & style)



💡 Use when: efficient domain style customization



## Guideline

Start with Prompts



Add RAG for missing knowledge



Fine-tune for domain behavior



# Contextualization Decision Tree

When to use Prompt Engineering, RAG, or Fine-tuning

Condition	Recommended Method	Examples
• Base LLM sufficient (with clear prompting)	 Prompt Engineering	ChatGPT Granola
• Missing or dynamic knowledge	 RAG (Retrieval-Augmented Generation)	Notion AI Notebook LM
• Need domain-specific behavior or style	 Fine-tuning	Bloomberg GPT



Guideline

Start with Prompts



Add RAG for missing knowledge



Fine-tune for domain behavior



# Transfer Learning — Big → Small

Large LLM → synthetic data → small model fine-tuning for cost and latency optimization



## 1) Large LLM

Frontier model as teacher

- Generate domain-specific synthetic Q/A, code, or labeled samples
- Apply quality filters and instructions



## 2) Synthetic Data

Curated training corpus

- Format as input → output pairs with rationales when useful
- De-duplication and balance across topics



## 3) Small Model + Fine-tune

Efficient deployment target

- Fine-tune smaller models (e.g., LLaMA variants, GPT nano-class)
- Use PEFT (LoRA/adapters) for faster, cheaper training

(\$) Lower cost & latency vs. large models

(⌚) Competitive quality on narrow tasks

(⚠) Validate data quality; monitor drift

Use cases: customer support assistants, coding copilots, niche Q&A. Start with prompts → add RAG if needed → fine-tune small models with synthetic data for cost-effective serving.



Section III

# Products, Tooling & Ops

---

Agents, low-code builders, MCP integrations, and evaluation practices





# AI Agents — Overview & Components

Definition, agent types, core components, and common platforms

## Definition

AI agents connect an LLM to tools/APIs and data so they can take actions autonomously or semi-autonomously to achieve goals.

### Workflow

Follows predefined steps or playbooks.

### Autonomous

Makes decisions based on context and goals.

## Platforms

Relay.app

ZPR

agents.ai

OpenAI agent kits

Example use case: weekly LinkedIn post summarizer that emails a digest.

## Core Components

### Reasoning

Plan, decide, and summarize.

### Memory

State, long-term notes, vector DBs.

### Tools

APIs, web actions, databases.

### Orchestration

Multi-step flows, task routing.

### Guardrails

Policies, constraints, safe tools.

Tip: Start simple with workflow agents, add memory and guardrails, then iterate toward autonomy as you gather feedback.



# Low-Code AI App Builders (Lovable)

Architecture flow and business value — democratizing app creation

## Architecture Flow



### User Prompt

Describe the app and goals in natural language.



### System Boost

Enhance with role, constraints, and style.



### LLM Plan

Generate file structure, APIs, and tasks.



### Agent Build/Test

Create code, run tests, debug iteratively.



### Deploy

App hosted and shareable instantly.



## Business Value

- Democratizes app creation for non-coders, expanding addressable market.
- Rapid prototyping reduces time-to-value for entrepreneurs and teams.
- Freemium + community templates drive viral distribution and retention.

Architecture applies broadly to low-code AI platforms; pair with evaluation and guardrails for production use.



Faster MVP cycles



Lower build costs



Broader creator base

Lovable example: users describe an app, the system enhances instructions, plans the build, agents generate and test code, then deploy a working app — often in minutes.



# Model Context Protocol (MCP)

Standardized layer for LLM  $\leftrightarrow$  services: Client • Server • Data Source

## MCP Architecture Stack



### MCP Client

Hosted on AI platform (e.g., ChatGPT, Claude)

Natural language requests



### MCP Server

Hosted by service provider (e.g., Zomato)

Maps intents to capabilities



### Data Source / Actions

Menu & catalog APIs

Orders & tracking

Payments

Goal: let LLMs call standardized tools without bespoke API knowledge. Scope capabilities explicitly; log and evaluate calls.



## Zomato Example



"Order a veg biryani from Zomato to {Fill in detailed address here}."

Client parses intent

→ Server maps to "place\_order"

APIs create order

## Benefits



Natural language interface



Faster integrations



Scoped, safer actions

Implementation tip: convert existing REST endpoints into MCP capabilities with clear schemas and permission checks.



# Evaluation — Quality, Safety, Reliability

Assess outputs, reduce risks, and enforce guardrails for production AI

## ⚠ Challenges

- Hallucination — confident but incorrect or unsupported claims.
- Bias — unfair, harmful or non-inclusive outputs.
- Non-determinism — variable answers to the same prompt.

Impact: reduced trust, compliance risk, unpredictable UX.



## 🔍 Methods

- LLM as judge — a stronger model scores relevance, correctness, tone.
- Automated metrics — task-specific checks (factuality, structure, toxicity).
- Manual review — sample audits to calibrate rubrics and catch edge cases.

Golden datasets

A/B comparisons

Score thresholds



## 🛡 Guardrails

- Prompt constraints — roles, style, and “don’t assume beyond context.”
- Tooling limits — allowlisted APIs, safe actions, timeouts.
- Red teaming — adversarial tests; feedback loops to retrain or refine.

Log prompts, context, outputs, and judgments for auditability.

Define task & rubric

Test on datasets

Track scores & regressions



# Toolkit + Key Insights

Tools by category • 10 takeaways for AI Product Managers

## AI PM Toolkit

### Discovery

ChatGPT Notebook Perplexity Mixpanel TextSQL

### Delivery

Jira Asana VZero Postbot Figma

### Distribution

Craftable Notion Genex

## Key Insights (10)

- 1 GenAI growth is driven by real utility, not just hype.
- 2 LLMs are probabilistic; engineer prompts and controls carefully.
- 3 RAG reduces context size and boosts factual accuracy.
- 4 Prompting is iterative—maintain a versioned prompt directory.
- 5 RLHF improves helpfulness, tone, and trustworthiness.
- 6 Agents extend LLMs with tools to automate real workflows.
- 7 Low-code builders democratize product creation at speed.
- 8 MCP standardizes tool access and accelerates integrations.
- 9 Rigorous evaluation (LLM-as-judge + metrics) is essential.
- 10 Win by blending tech depth with user empathy and business sense.

Tip: Start lightweight—use prompts and RAG first; fine-tune or build agents only when the use case clearly demands it.



# Glossary

Key AI/ML terms used in this masterclass

Term	Definition
LLM	Large Language Model — neural network trained to predict the next token, enabling language understanding and generation.
Token	Numeric unit of text used by models; pricing and context limits are measured in tokens.
Transformer	Architecture (2017) that processes sequences in parallel using attention mechanisms.
Attention	Mechanism that weighs relationships among tokens to capture context.
RAG	Retrieval-Augmented Generation — combines vector search of relevant docs with LLM generation.
Embedding	Vector representation of text that captures semantic meaning for similarity search.
Fine-tuning	Adapting a pre-trained model to a domain/task using labeled examples.
PEFT	Parameter-Efficient Fine-Tuning (e.g., LoRA, Adapters) — updates a small subset of weights.



# Closing — Build • Measure • Learn

Action + feedback + reflection → mastery through building



## Build. Measure. Learn.

Ship small, gather signal, reflect deeply, then iterate.



### Build

Create small MVPs; ship weekly.



### Feedback

User signals + LLM-as-judge evals.



### Reflection

Retros; refine prompts/RAG.



### Mastery

Compound learning via iteration.

➤ Start with Prompts



➤ Add RAG for missing knowledge



➤ Fine-tune for domain behavior



Prioritize user value • evaluate rigorously • iterate fast

Thank you! Questions?